



# Data Ingestion with Delta Lake

---





# Agenda

- Introduction to Delta Lake
  - Overview of Delta Lake
  - Benefits of using Delta Lake
- Delta Lake Architecture
  - Components of Delta Lake
  - How Delta Lake works
- Comparison with Traditional Data Lakes
  - Differences between Delta Lake and traditional data lakes
  - Advantages of Delta Lake
- Creating and Loading Delta Tables
- Basic Data Transformations
- Hands-On Labs and Demos



# Introduction to Databricks

---

## Unified Data Platform

Databricks provides a cohesive environment for data processing and machine learning on Apache Spark.

## Collaborative Workflows

It enables better collaboration among data engineers, scientists, and analysts by streamlining workflows.

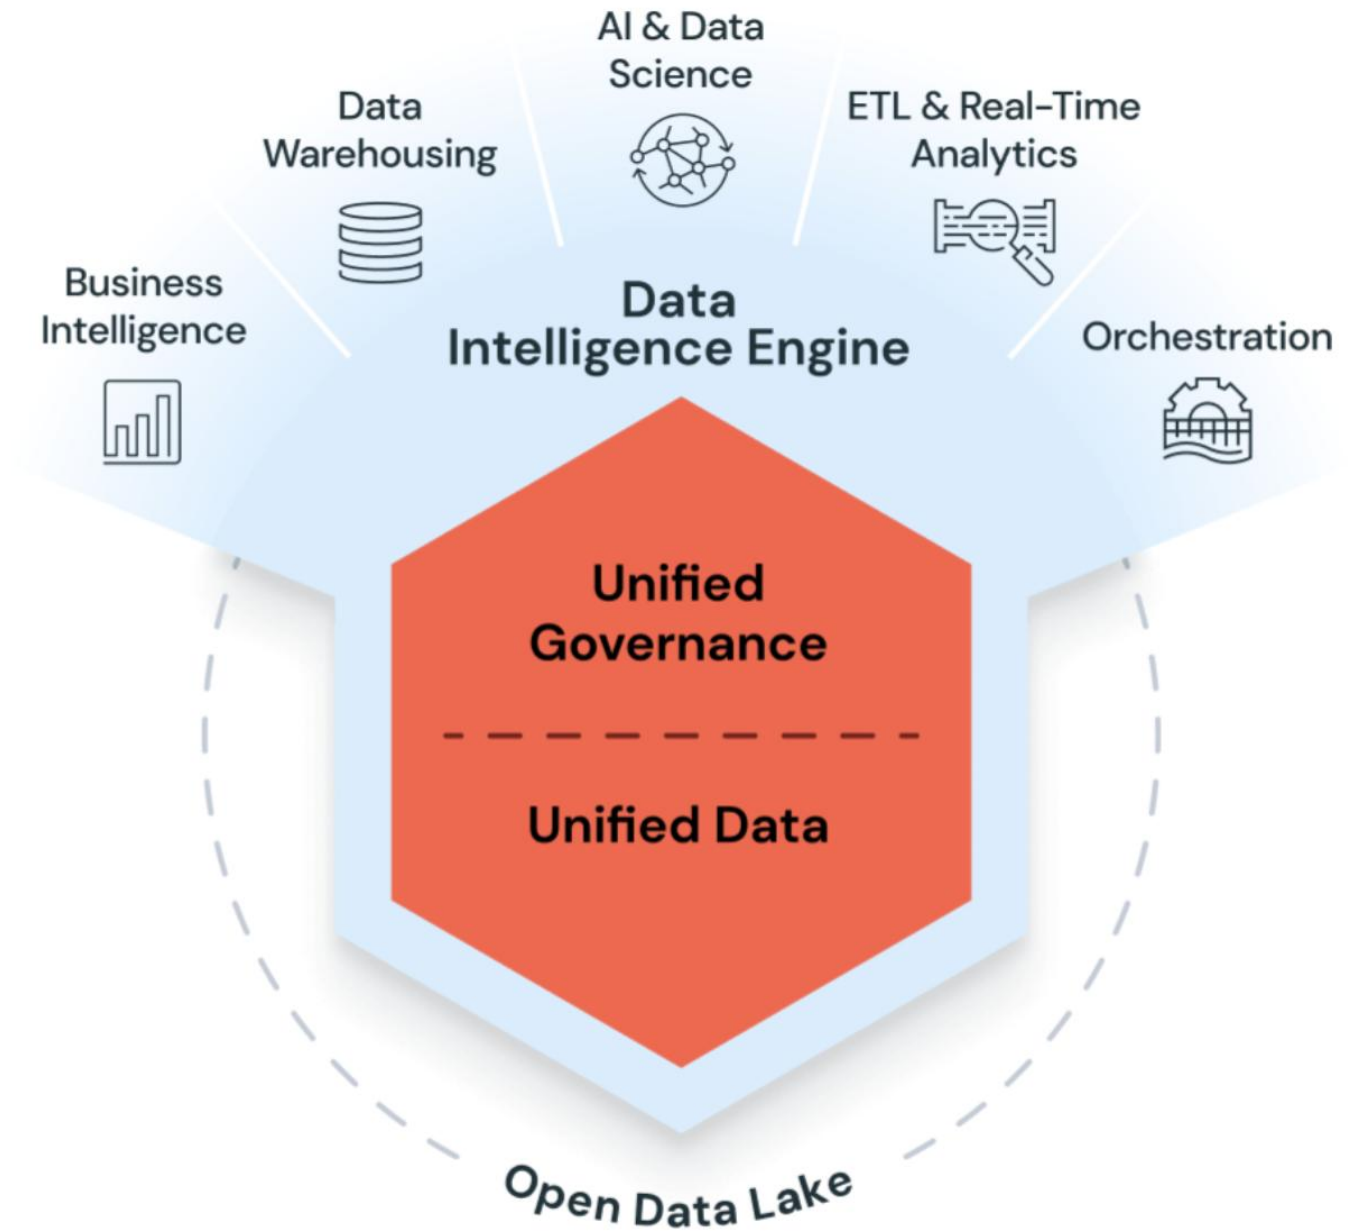
## Batch and Real-Time Analytics

Databricks supports both batch and real-time analytics, making it adaptable for various data-driven applications.

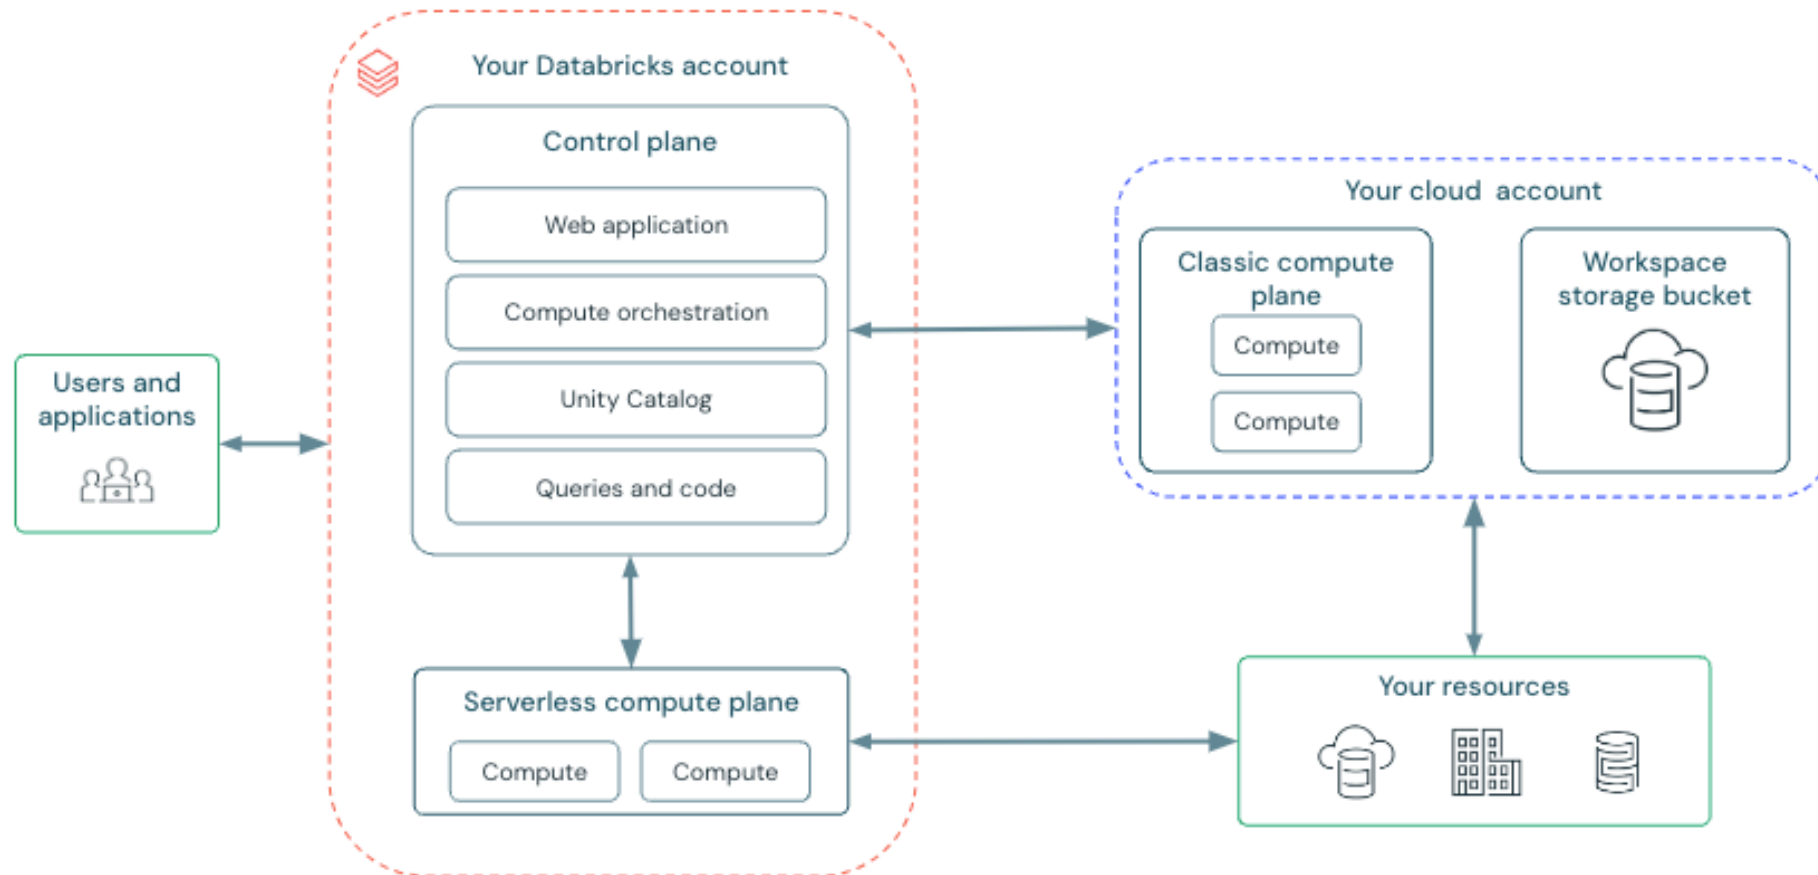
## Integration with Tools

It integrates seamlessly with various data sources and tools to enhance data operations.

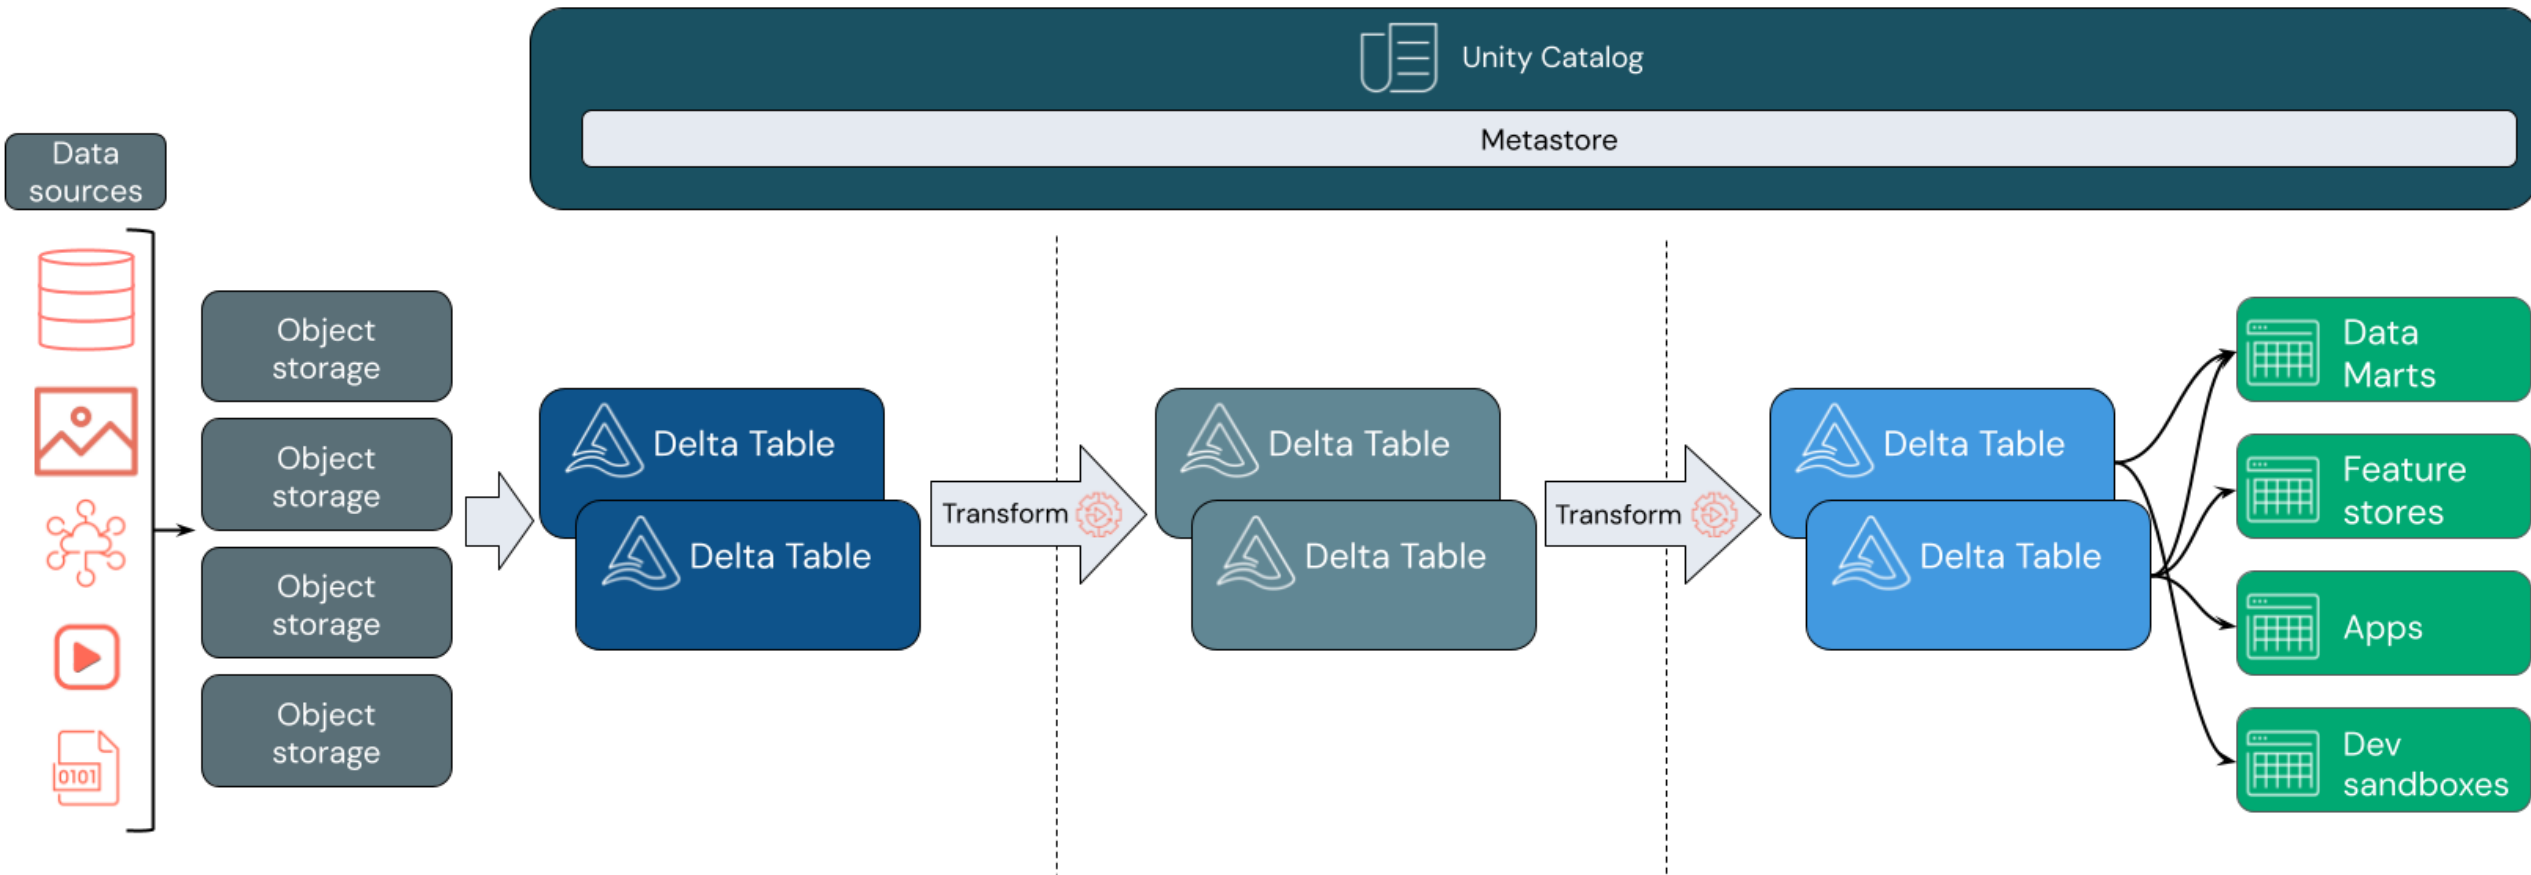
# Databricks



# Databricks High Level Architecture



# Lakehouse



# What is Delta Lake?



Open-source storage layer

Designed for data lakes



ACID transactions

Enhances Apache Spark  
Optimizes big data workloads



Schema enforcement and evolution

Ensures data integrity  
Adapts to changes



Scalable and reliable

Supports large-scale data  
Ensures high performance

# Delta Lake Architecture Overview



## Foundation of Delta Lake Architecture

Built on Apache Spark and Parquet files



## Transaction Log

Tracks all changes using .json and checkpoint files



## Time Travel

Enables data versioning



## Data Processing

Supports both batch and streaming data processing



# Key Differences

---



## Data Consistency and Schema Enforcement

Traditional data lakes often struggle with these issues

Delta Lake solves them with ACID transactions



## Time Travel

Allows easy rollback of data changes  
Facilitates auditing of data changes

# Comparison Table




Feature	Traditional Data Lakes	Delta Lake
Data Consistency	Limited (eventual)	ACID transactions
Schema Enforcement	Manual & error-prone	Built-in schema enforcement
Data Versioning	Usually not available	Built-in time travel
Performance Optimizations	Limited	Z-ordering, compaction

# Loading Data into Delta Tables



Ingest data from  
various file formats

CSV files  
JSON files  
Parquet files



Batch ingestion  
method

Use  
`write.format("delta").save()`  
for batch ingestion



Streaming ingestion

Overview provided  
Detailed coverage later

# Creating Delta Tables

- Syntax to create Delta tables from existing data
  - Steps to convert existing data into Delta format
  - Commands and examples for implementation
- Managed vs external Delta tables
  - Differences between managed and external tables
  - Advantages and disadvantages of each type
- Creating tables using SQL and PySpark
  - SQL commands for table creation
  - PySpark methods for table creation



# DataFrame Transformations

---

- Select, filter, add columns, drop columns
  - Using PySpark DataFrame API
  - Example: Add a total price column (Quantity \* UnitPrice)

# SQL Transformations



Query Delta tables  
using Spark SQL

Utilize Spark SQL to  
interact with Delta  
tables



Simple SELECT,  
WHERE, and  
aggregate queries

Perform basic SQL  
operations  
Filter data using WHERE  
clause  
Aggregate data for  
summary statistics



Example: Total sales  
per customer

Calculate total sales for  
each customer  
Demonstrate practical  
application of SQL  
queries

# Hands-On Labs and Demos

- Create Delta Table from provided sales CSV
  - Import sales data from CSV file
  - Initialize Delta Table
- Load sales data into Delta Table
  - Insert data into Delta Table
  - Verify data loading
- Perform simple DataFrame transformations
  - Apply basic transformations
  - Filter and aggregate data
- Run SQL queries on Delta Table
  - Execute SQL commands
  - Analyze query results